

Forecasts


#	Forecast	Prob	Resolves via	Tag (B - base case A - alternative world)
1	Best open-weight model within 3mo / <50 Elo of frontier, end-2028	0.80	LMarena/ Epoch	B — diffusion
2	No system reaches human-expert parity on 32h+ ML-R&D (METR RE-bench ≥ 1.0 sustained) by end-2028	0.70	METR RE-bench	A — RSI master gauge
3	Open-weight share of tokens $\geq 40\%$ by end-2027 (from 33% in 2025)	0.65	OpenRouter /a16z	B — diffusion
4	Blended GPT-4-class API price <\$1/M tokens by end-2026	0.70	Artificial Analysis price index	B — value leaves model layer
5	A top-5 frontier lab ceases to be an independent model-API business by end-2028	0.55	filings	B — value migrates up-stack
6	China does not reach yield-viable EUV $\leq 5\text{nm}$ volume by end-2028	0.80	TechInsights	B — chips bind China
7	China's annual net power adds >5 $\times$ the US's in 2027	0.85	China NEA/EIA	B — energy asymmetry
8	$\geq 25\%$ of announced US 2026-27 data-center capacity delayed/relocated by end-2028	0.70	Data Center Watch/LBNL	B — US power bind
9	US chip controls degrade to transactional (revenue-share / case-by-case) by end-2028	0.65	BIS/NVIDIA filings	B — controls leak (COCOM)
10	No full kinetic PRC invasion of Taiwan by end-2028	0.85	events	A — fracture
11	A major AI-buildout financial-stress event (top-10 infra default/impairment or a big-4 cutting AI-capex guidance YoY) by end-2027	0.35	filings/market data	A — Perez bust
12	US white-collar unemployment attributable to AI stays <2pp above trend through 2027	0.65	BLS	A — societal-upheaval
13	USD share of allocated FX reserves stays >50% thru 2028	0.80	IMF COFER	A — Dalio/macro
14	A Chinese-origin open-weight family (Qwen/ DeepSeek lineage) is #1 by global derivative-model count by 2027	0.70	HuggingFace	B — tests China's deployment win
15	≥ 1 top-3 US AI lab sites a ≥ 1 GW cluster outside the US (Gulf/allied) primarily for power access, by 31 Dec 2028	0.55	press/filings	A/B — tests the US energy bind

Framework & Holistic Synthesis


My argument has three parts:

1. **Diffusion beats concentration:** frontier AI progress diffuses faster than it can concentrate - even in the case of a self-improving superintelligence.
2. **Value migrates from model layer:** as the model layer commoditizes value migrates in a barbell - downstream to the energy/hardware layer and upstream to distribution/services.
3. **China wins battle of mercantilist constraints:** chips and EUV bind China today and these matter most for training, but that lead diffuses/commoditizes. Instead, the binding constraint migrates to inference and energy where the US is disadvantaged allowing China to win the deployment tier.

Diffusion beats concentration:




Source: METR, Measuring AI Ability to Complete Long Tasks (2025)




Source: best-closed – best-open (LMSYS / LMArena Chatbot Arena)

The key race is between AI capability improvement and the diffusion and commoditization of that intelligence. Capability is rising fast — METR's autonomous task-horizon at 50% reliability has gone from ~2 seconds to ~2.3 hours since 2019. But diffusion is faster: the Elo gap to the best open-weight model has collapsed from 250 in 2023 to 48 in 2026, and the lag from a closed state-of-the-art release to open-weight parity has fallen from ~24 months (GPT-3 in 2020 → OPT/BLOOM in 2022) to ~3 months by 2025 (Epoch AI). This is structural, not incidental, because (1) distillation is a one-way ratchet — followers copy an expensive teacher model at near-zero marginal cost — and (2) algorithmic-efficiency gains are non-excludable, leaking through publication, the internet, and a global labour market. The base rate for technological leads is that they are transient: the US held an atomic monopoly for just four years (1945–1949) despite total secrecy, an ocean of separation, and a physical-materials chokepoint tighter than anything in software. Even a self-improving Bostrom superintelligence would keep an unassailable lead only if (1) takeoff outpaced diffusion — hard when algorithms leak through the internet and labour mobility (the Klaus Fuchs channel) — and (2) the cognitive lead converted into a decisive resource/compute advantage, which is unlikely while the AI remains human-controlled: a US lab with a superintelligence is unlikely to kill everyone in China or seize all the world's compute for itself. So I forecast the diffusion side wins — open weights stay within ~3 months / <50 Elo of the frontier through 2028 (F1, 80%); no system crosses the recursive-self-improvement threshold that would let a lead compound, i.e. human-expert parity on 32h+ ML-R&D tasks, by 2028 (F2, 70%); and open weights reach ≥40% of tokens by 2027 (F3, 65%).

Value migrates from the model layer:



Source: blended GPT-4-class API list prices (OpenAI, TokenCost/BenchLM)




Source: OpenRouter / a16z 100T-token study

Because diffusion outpaces capability, rent has drained from the model layer: the cost to reach a fixed (GPT-4-class) capability has fallen ~98% from 2020 to 2026, while open-weight models have taken a rising share of tokens — from 4% in 2023 to 33% in 2025 (F4: blended GPT-4-class API price falls below \$1/M tokens by end-2026, 70%). This open-base commoditization has precedent — MySQL/PostgreSQL erasing Oracle's database rent; RISC-V eroding ARM's licensing chokepoint. In each case the clever core commoditized and value moved off it in two directions. Up-stack, into distribution and services: Databricks built managed services on open-source Spark and thrived, while Cloudera — pure Hadoop distribution — failed; open-source Android protects Google's search, ads, and Play services ("commoditize your complement"). Down-stack, into the physical layer: TCP/IP commoditized and value moved into fibre and data centres; electrification's rent moved from the dynamo to the grid, utilities, and equipment makers (GE/Westinghouse); Standard Oil won by owning pipelines and rebates, not by refining better. I expect the same barbell here. Value accrues up-stack to the hyperscalers that own distribution and services (cloud plus consumer surfaces like Google Drive), and down-stack to whoever owns the proprietary physical standard — NVIDIA/CUDA, the Wintel of AI (the IBM-PC-clone pattern, in which the assembly commoditized and value concentrated in the proprietary standard), already at ~38x its 2019 data-centre revenue at ~73% gross margins; and TSMC, whose advanced-packaging capacity has climbed from 8k to 130k wafers/month between 2022 and 2026. The corollary is a forecast that the model layer stops being a standalone business: a top-5 frontier lab ceases to operate as an independent model-API company — absorbed, acquired, or pivoted — by 2028 (F5, 55%).

China wins the battle of mercantilist constraints:

As rent leaks from the model layer, the binding constraints become mercantilist. China's constraint is chips: NVIDIA's China revenue fell from 26% (FY2022) to 13% (FY2025); the BIS Entity List grew ~5.5x to ~1,300 entries between 2016 and 2025; and China cannot easily make its own leading edge — the Netherlands' ASML holds ~100% of sub-5nm EUV lithography, and SMIC's rapid indigenous DUV-only progress (28nm in 2015 to ~5nm-class in 2025) stalls at uneconomic ~30–40% yields, with no path further without EUV. But I forecast this gate does not bind China, for two reasons. First, the controls are porous — the COCOM/Wassenaar lesson is that export controls delay but do not deny, and an estimated 660k+ H100-equivalents reached China in 2025 anyway (F9: US controls degrade to transactional — revenue-share or case-by-case licensing, or a net loosening — by 2028, 65%). Second, and more fundamentally, chips bind training, which sets the model frontier — and that frontier diffuses. Chips do not bind inference, where most of the economic value will

accrue (F6: China does not reach yield-viable EUV $\leq 5\text{nm}$ volume production by 2028, 80%).



The US's constraint is energy. The interconnection queue has grown ~5x since 2015 (to ~2,600 GW), and large-transformer / high-voltage lead times have widened to ~28–40 months, against data-centre demand rising from ~35 GW in 2025 to an estimated ~80 GW by 2030 — and a hardening political backlash: Gallup finds 71% of Americans oppose a data centre near them, understandable when PJM's independent market monitor attributed ~63% of the 2025/26 capacity-cost increase to data centres. China, by contrast, added ~540 GW of net capacity in 2025 versus the US's ~56 GW — and even that US figure was mostly solar, the wrong kind of power for 24/7 data-centre load (F7: China's annual net power additions exceed 5x the US's in 2027, 85%; F8: $\geq 25\%$ of announced US 2026–27 data-centre capacity is delayed, cancelled, or relocated by 2028, 70%). Energy has been decisive before — coal underwrote Britain's 18th–19th-century industrial revolution; oil abundance underwrote US hegemony in the 20th — but only when coupled with the absorptive capacity to convert it (the petrostate resource curse, and the energy-rich Soviet Union that still lost, are the counter-examples). China uniquely has both.

Liebig's law holds that growth is set not by total resources but by the scarcest input. The core claim of this framework is that the binding minimum migrates — from China's EUV/chip constraint, which governs the fast-commoditizing training frontier, to the US's energy constraint, which governs the deployment phase where the durable value lives.

Inversion

The page limit forces a focus on the base case (~55%), but tail cases worth considering:

- **Perez-style demand/ROI bust (~20%):** as in the 2000–02 telecom/fibre crash, the crash is the turning point between the financial-capital-led installation frenzy and the deployment phase — a shakeout resolved by government industrial policy and re-regulation, plausibly against a late-stage Dalio debt cycle in which monetization and financial repression fund the backstop and accelerate dollar erosion (F11: a major AI-buildout financial-stress event by 2027, 35%; F13: the dollar holds $>50\%$ of allocated reserves through 2028, 80%).
- **TSMC/Taiwan fracture (~13%):** a military action that throttles both US and Chinese progress at once by cutting leading-edge supply (F10: no full kinetic invasion of Taiwan by 2028, 85%).
- **Superintelligence takeover (~12%):** a runaway AI, possibly outside human control — the one branch in which a lead does lock in (F2 is the gauge that would signal we are entering it).
- **Societal upheaval (cross-cutting):** AI-driven labour displacement on the scale of the industrial revolution's destruction of agricultural work, but faster — a risk that can accompany any branch (F12: AI-attributable white-collar unemployment stays $<2\text{pp}$ above trend through 2027, 65%).
- **Orbital data-centers:** SpaceX-style orbital data-centers could remove US's territorial energy constraint.

Analytical Appendix

AI has the potential to be a unique event in human history making it difficult to analyse. The approach I have taken has a number of components:

- **Scaled qualitative analysis:** Similar to how computers allowed the scaling of quantitative analysis I have used agentic workflows of LLMs to scale qualitative analysis.
- **Outside view first:** Tetlock in his superforecasting research has shown that (borrowing Isaiah Berlin's dichotomy) hedgehogs who interpret everything through one world model tend to underperform foxes who hold multiple conflicting models simultaneously. This is similar to Charlie Munger's idea of using a lattice work of diverse mental models that can outperform experts even in their own domains. Thus the pipeline starts by generating 100s of reference examples first before developing any theory or conclusions.
- **Disconfirming evidence:** Although less relevant for this point-in-time analysis, every generated note includes a steel-man of the other side and key metrics to track that would indicate a move away from the base-case into one of the tail-cases for example. Thus the model can be run continuously and evolve over time.
- **Pipeline construction:** Initial versions of my pipeline failed because
 - **Combinatorial explosion:** Initial versions included a combinatorial explosion of every possible factor with every other factor, but what worked better was to use a multi-lens completeness critic to mine for other angles.
 - **Flattening too early:** I looked to generate cruxes too early resulting in a flattening of key interactions e.g. physical inputs vs intelligence collapsed a vector of chips, energy, grid, material, labour and capital into one axis with no nuance for different actors.
 - **Status quo predictions:** A lot of the predictions were just independent projections of the status quo, grounding each topic in graphs that best represented each topic helped create a trend to predict off grounded in the historical examples.
 - **No interactions:** In order to generate cross-topic interactions I 1. look for coupled loops and cycles over the 42 merged variables 2. ground in theory to prevent hypothesising over every edge of the graph (particularly Dalio, Perez, Bostrom and Acemoglu) 3. prioritise to determine which loop/constraint dominates (Liebig) 4. Tie the loop to a historical instance and a live gauge
- **Future work:** Some things I would like to be able to do in future work include:
 - **Human-in-the-loop:** The process for generating this analysis was highly hand-made which does not scale. I would like to leverage more Karpathy type recursive improvement on the pipeline itself (although I was constrained by my token budget here)
 - **Evaluation:** I also wasn't able to build robust evaluation workflows to both validate the notes and data I used, nor the analysis that was built on it.
 - **Run model live:** Update the model to look for disconfirming evidence (see disconfirming evidence table below) live and update my predictions accordingly

Global variables:

My analysis generated 9 core topics within which I have 42 key global variables

1. Macro-financial loop:
 - a. GV7 Energy-cost trajectory (abundance vs J-curve)
 - b. GV36 Central-bank independence vs fiscal-dominance
 - c. GV37 Dollar reserve status / marginal UST demand ("who buys the bonds")
 - d. GV38 Treasury-market plumbing / non-linear stress capacity
 - e. GV39 Neutral real rate (r^*) / global loanable-funds balance
 - f. GV40 AI productivity offset (macro TFP $\times$ timing)
 - g. GV41 Inflation transmission / distributional channel
 - h. GV42 Political resolution path (who bears the adjustment)

2. Scaling and the model layer:
 - a. GV9 Scaling axis & data wall (where capability originates)
 - b. GV10 Efficiency compression & open-weight commoditization (the diffusion clock)
 - c. GV12 Frontier AI talent concentration & mobility
 - d. GV15 Value-capture locus (where durable rent accrues)
3. Agentic diffusion and labour
 - a. GV13 Diffusion / deployment velocity (capability-to-deployment gap)
 - b. GV17 Labor reinstatement ratio (task creation vs displacement)
 - c. GV18 Wage-appropriability & aggregate-demand feedback
 - d. GV32 Political backlash channel, coalition speed & national-security override
4. Trade order:
 - a. GV19 US-China power-transition trajectory
 - b. GV20 Security-fragmentation / bloc structure (weaponized interdependence)
 - c. GV23 Supply-chain integration floor (stickiness, agglomeration)
 - d. GV24 Trade-order anchor / rule-writer
 - e. GV25 Institutional adaptation capacity (WTO / dispute settlement)
 - f. GV29 Managed-trade fiscal & state-capacity sustainability
5. Tech-control paradigm
 - a. GV1 Leading-edge logic fabrication (sub-5nm capacity & indigenization)
 - b. GV11 Decisive-capability threshold & its timing (the master clock)
6. AI buildout
 - a. GV2 EUV / EDA tooling chokepoint (the gate behind the gate)
 - b. GV3 Advanced packaging (CoWoS) + HBM
 - c. GV6 Power & grid execution (electrons, transformers, interconnection)
 - d. GV8 Critical-minerals & grid-equipment refining (China counter-lever)
 - e. GV14 Demand realization / killer-app ROI
 - f. GV16 AI-buildout capital structure & financial fragility
7. Regulatory/standards fragmentation
 - a. GV5 Software/platform moat (CUDA / interconnect / EDA-software)
 - b. GV30 AI regulation, liability & standards regime
 - c. GV31 Training-data legality / IP regime
 - d. GV33 Antitrust / structural intervention capacity
8. Multipolar bloc formation
 - a. GV4 Chip-access export-control regime (BIS controls + coalition compliance)
 - b. GV26 Swing-state / connector alignment & Global-South standards contest
 - c. GV27 Great-power auction / second-pole credibility (China as outside option)
 - d. GV28 Gulf neutrality-premium / dual-stack hub viability
9. Tail discontinuities
 - a. GV21 Taiwan/TSMC discontinuity & Beijing's action calculus
 - b. GV22 Great-power war scope, nuclear firebreak & exogenous chokepoint shocks
 - c. GV34 State-capture / control structure of decisive AI advantage
 - d. GV35 Safety-catastrophe clampdown regime

Feedback loops:

My analysis surfaced a few key reinforcing loops:

- **Recursive superintelligence:** capability (GV11) → commoditization (GV10) → scaling/data (GV9) The only reference-class-free loop; if it self-closes the field freezes and one lab wins — but every gauge reads sub-unity, so it's dormant. Bostrom's argument builds on three assumptions that takeoff is fast, compute-unconstrained and non-diffusing. One example of plateau risk is AI Winter II where expert-systems (1984-1993) hit a reasoning wall.

- **Compute–capex flywheel:** capability (GV11) → value capture (GV15) → capital (GV16) → chips/fabs (GV1) . The loop moving the trillions — capability earns rent, funds capex, buys compute, lifts capability. Braked only by demand/ROI. A historical example is the ~\$1T of 1996-2001 fiber buildout where the demand curve arrived late and prices collapsed 99%+.
- **Diffusion/commoditization:** commoditization (GV10) → deployment (GV13) → demand (GV14) → capability (GV11). Cheap open models spread, drive usage and training signal back into capability — and erase every lead in the process (the diffusion clock). A historical example is MySQL/PostgreSQL & LAMP stack vs Oracle/DB2 where the commodity base wins the cost-sensitive majority and rent migrates to hosting/distribution. Another example is nuclear weapon proliferation where the Klaus Fuchs/Theodore Hall espionage network compressed the Soviet timeline directly.
- **Chip-chokepoint moat:** chips (GV1) → capability (GV11) → China-as-pole (GV27) → export controls (GV4) US controls protect the chip lead, keep China behind, and justify more controls — self-reinforcing bloc formation that leaks over time. A historical example is VoIP defeating PTT hardware regulation and RISC-V vs ARM is a goods-based control regime that is structurally defeated by cloud-API/inference leg and by adversarially open-sourcing.
- **Talent flywheel:** talent (GV12) → capability (GV11) → value capture (GV15). Elite talent concentrates → better models → more rent → attracts more talent; throttled only by US visa policy. Examples include Taiwan reverse migration back to TSMC although this failed in post-Soviet Russia's case.
- **Deployment/implementation not exploration:** deployment velocity (GV13) → adoption/demand (GV14) → proprietary usage-data & value capture (GV15) → reinvestment → more deployment (GV13) This is essentially Kai Fu Lee's AI Superpowers argument where commoditized models favour China as it has 1. data advantage 2. gladiatorial entrepreneurial culture 3. techno-utilitarian government 4. manufacturing ecosystem particularly for physical AI with robots and drones (China has ~50-60% of industrial robot installs). Note that even if China captures this end of the barbell, the US will still capture rent in services and distribution layers and potentially around standards like NVIDIA's CUDA

And a few key balancing/braking loops

- **Demand/ROI fragility ceiling:** capital (GV16) → buildout/chips (GV1) → demand/ROI (GV14) The buildout races ahead of revenue; if ROI never validates, financing snaps — the Perez-bust brake. Examples include British Railway mania (1844-47) with ~7% GDP/year ending in 1866 Overend Gurney collapse.
- **Power/grid ceiling:** demand (GV11) → power & grid (GV6) → deployment (GV13). Electrons, transformers, and interconnection queues cap how fast compute can actually be deployed — the binding US constraint. Example is Alcoa's aluminium energy-geography lock-in (1886-1945) where compute will migrate to where the cheap firm power already is.
- **Minerals MAD:** export controls (GV4) → critical minerals (GV8) → US power buildout (GV6) US chip controls invite China's minerals retaliation, which taxes the US buildout — mutual assured disruption. Example is 2010 China rare-earth embargo on Japan (Senkaku) showing effectiveness floored by WW2 US synthetic rubber crash program (1942-45) showing how fast denial can be beaten.
- **Fiscal-dominance / bond-market loop:** real rate r^* (GV39) → dollar/UST demand (GV37) → central-bank independence (GV36) → who-bears-adjustment (GV42) Financing the buildout strains the bond market → repression/fiscal dominance → dollar erosion — the macro ceiling. Examples include 1844-47 railway mania -> bank charter act suspension -> 1890s Barings crisis and a capital export mania meeting a monetary boundary.
- **Labor / underconsumption brake:** capability (GV11) → labor displacement (GV17) → wages & demand (GV18) → demand/ROI (GV14) Deployment displaces labor, cuts wages and aggregate demand, and undercuts the very demand the buildout needs — the Layer-3 feedback. Acemoglu-Restrepo task-displacement runs ahead of reinstatement.

Liebig sequencing and durability analysis

Given the lattice of mercantilist constraints the obvious question is which one would matter most and when. Liebig sequencing argues that binding minimum can change and history can be a guide as to how durable each is (from most to least)

- **Standards/networks:** Durable. examples include AC-power, railway gauge, Wintel, dollar/SWIFT, Western Union. Implications include CUDA lock-in, hyperscaler-cloud aggregation, app-layer aggregation.
- **Energy & Absorptive capacity:** Durable when coupled. Coal in Britain, Oil in the US but petrostates and Soviet Union are counter examples of how they also need absorptive capacity - so petrostates may again find it hard to exploit their energy advantage.
- **Precision tooling:** Durable but rare. examples include Zeiss/ASML which beat Nikon/Canon.
- **Capital/finance:** Durable but insufficient.
- **Manufacturing capacity:** Durable for a generation. DRAM went from the US to Japan to Korea to China. Similar stories with steel, shipbuilding and solar.
- **Talent/knowledge:** Diffuses. E.g. Manhattan project to USSR 49, same with Operation Paperclip. Especially as so much AI talent in the US is Chinese and therefore could repatriate.
- **Materials/minerals:** Can be routed around e.g. nitrate with Haber-Bosch or rubber to synthetics or 2010 REE to Lynas.
- **Inventions:** Not durable. Often commoditized. Examples include electrification (1880-1930) where dynamo/AC induction motor commoditized and value landed on the grid. Similarly with the telegraph in the 1840s-1990s where Morse's patent and telegraph key were worked around and licensed away, the value landed in Western Union's network and transatlantic cable layers. Also the case with steam where the value moved from Watt's engine to coupling coal with skilled labour.

Superintelligence has no base case

The one example that does not have a historical base case is a self-improving superintelligence (although even there I draw on examples to argue for diffusion over concentration). Nonetheless, I use Bostrom's framework to frame two questions

1. **How a superintelligence could gain an unassailable lead:** 1. Initial cognitive lead converted into resource/compute lead 2. Initial cognitive lead optimized into a time lead where optimisation dominates recalcitrance 3. Initial cognitive lead leads to wealth/financial resources. The two constraints are physical constraints and diffusion.
2. **What a superintelligence could do with that lead:** There are two regimes: 1. human-led 2. superintelligence led. The former at least would likely be bound by some limitations e.g. the US would probably not tolerate killing everyone in China although as in the recent Iran conflict it might try to bomb all AI facilities and energy infrastructure and potentially the superintelligence could be used to hack China's nuclear weapons etc. For the latter it is much harder to reason about, but potentially there could be no limitations on what it might do including cyber, bio, persuasion, compute seizure, market manipulation.

Disconfirming evidence

It is important that I know up front the key thresholds that if breached would make me change my mind. These are defined in the table below

Base-case bet	Kill signal (measurable)	Breaks if...	Flips to	Watch / source	Forecast
Diffusion beats concentration	Closed-open capability gap	best-closed – best-open Elo gap rises for ≥2 consecutive quarters(reverses	Concentration / lead compounds	LMARena; Epoch	F1, F3

	<i>re-widens</i> instead of closing	250→48), or open-weight parity lag climbs back >6 months			
No recursive-self-improvement takeoff	AI crosses human-expert parity on long-horizon R&D	METR RE-Bench ≥ 1.0 at the 32-hour budget, sustained	RSI takeover branch	METR RE-Bench	F2
Value migrates off the model layer	Model-layer rent <i>holds</i> instead of commoditizing	a top frontier lab sustains >50% gross margin on model-API revenue for ≥ 4 quarters while growing, or fixed-capability API price flat for 12 months	Model moat holds (concentration)	lab filings; Artificial Analysis price index	F4, F5
Chips don't bind China (they bind training, not inference)	China starved of inference-scale compute	Chinese-origin open-weight token/derivative share falls ≥ 2 quarters, or Huawei/SMIC accelerator output collapses under tightened controls	US contains China (controls work)	HuggingFace; OpenRouter; TechInsights	F14, F6
The US binding constraint is energy	US clears the power gate faster than expected	US firm (dispatchable) additions >90 GW/yr, transformer lead times fall <18 months, or US data-center delay rate <10%	US not energy-bound → US wins both tiers	EIA; LBNL "Queued Up"; Data Center Watch	F8
No demand/ROI bust (base, not Perez tail)	AI-buildout financing cracks	a top-10 AI-infra borrower defaults/impairs, a big-4 cuts YoY AI-capex guidance, or a $\geq 30\%$ drawdown in the AI-datacenter equity complex	Perez bust branch (~20%)	filings; market data	F11
No Taiwan fracture	PRC moves from gray-zone to kinetic	declared blockade/quarantine with kinetic engagement, or amphibious action	Taiwan fracture branch (~13%)	events; PLA ADIZ activity	F10
The dollar/monetary order holds (no Dalio break)	Reserve status or Treasury plumbing cracks	USD <50% of allocated reserves, or a forced Fed backstop / failed-auction plumbing event	Dalio/fiscal-dominance branch	IMF COFER; Treasury market	F13
No fast societal rupture	Labour displacement outruns reinstatement	US professional/management unemployment >2pp above 2015-24 trend	Societal-upheaval (cross-cutting)	BLS	F12